



```
In [60]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as st
import statsmodels.api as sm
from scipy.stats import linregress
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: df = pd.read_csv('yellow_tripdata_2020-01.csv')
df.head()
```

```
Out[2]:   VendorID tpep_pickup_datetime tpep_dropoff_datetime passenger_count trip_distance RatecodeID store_and_fwd_flag PULocationID DOLocationID payment_type fare_amount extra mta_tax tip_amount tolls_amount improvement_surcharge total_amount congestion_surcharge
0      1.0 2020-01-01 00:28:15 2020-01-01 00:33:03           1.0          0.000000        1.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000
1      1.0 2020-01-01 00:35:39 2020-01-01 00:43:04           1.0          0.000000        1.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000
2      1.0 2020-01-01 00:47:41 2020-01-01 00:53:52           1.0          0.000000        1.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000
3      1.0 2020-01-01 00:55:23 2020-01-01 01:00:14           1.0          0.000000        1.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000
4      2.0 2020-01-01 00:01:58 2020-01-01 00:04:16           1.0          0.000000        1.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000          0.000000
```

```
In [3]: df.shape
```

```
Out[3]: (6405008, 18)
```

```
In [4]: df.dtypes
```

```
Out[4]: VendorID          float64
tpep_pickup_datetime    object
tpep_dropoff_datetime    object
passenger_count          float64
trip_distance            float64
RatecodeID               float64
store_and_fwd_flag       object
PULocationID             int64
DOLocationID             int64
payment_type              float64
fare_amount                float64
extra                      float64
mta_tax                     float64
tip_amount                  float64
tolls_amount                 float64
improvement_surcharge      float64
total_amount                  float64
congestion_surcharge       float64
dtype: object
```

```
In [5]: df['tpep_pickup_datetime'] = pd.to_datetime(df['tpep_pickup_datetime'])
df['tpep_dropoff_datetime'] = pd.to_datetime(df['tpep_dropoff_datetime'])
```

```
In [6]: df['duration'] = df['tpep_dropoff_datetime'] - df['tpep_pickup_datetime']
df['duration'] = df['duration'].dt.total_seconds()/60
df
```

```
Out[6]:
```

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count
0	1.0	2020-01-01 00:28:15	2020-01-01 00:33:03	
1	1.0	2020-01-01 00:35:39	2020-01-01 00:43:04	
2	1.0	2020-01-01 00:47:41	2020-01-01 00:53:52	
3	1.0	2020-01-01 00:55:23	2020-01-01 01:00:14	
4	2.0	2020-01-01 00:01:58	2020-01-01 00:04:16	
...
6405003	NaN	2020-01-31 22:51:00	2020-01-31 23:22:00	
6405004	NaN	2020-01-31 22:10:00	2020-01-31 23:26:00	
6405005	NaN	2020-01-31 22:50:07	2020-01-31 23:17:57	
6405006	NaN	2020-01-31 22:25:53	2020-01-31 22:48:32	
6405007	NaN	2020-01-31 22:44:00	2020-01-31 23:06:00	

6405008 rows × 19 columns

```
In [7]: df = df[['passenger_count','payment_type','fare_amount','trip_distance','duration']]
df
```

```
Out[7]:
```

	passenger_count	payment_type	fare_amount	trip_distance	duration
0	1.0	1.0	6.00	1.20	4.800000
1	1.0	1.0	7.00	1.20	7.416667
2	1.0	1.0	6.00	0.60	6.183333
3	1.0	1.0	5.50	0.80	4.850000
4	1.0	2.0	3.50	0.00	2.300000
...
6405003	NaN	NaN	17.59	3.24	31.000000
6405004	NaN	NaN	46.67	22.13	76.000000
6405005	NaN	NaN	48.85	10.51	27.833333
6405006	NaN	NaN	27.17	5.49	22.650000
6405007	NaN	NaN	54.56	11.60	22.000000

6405008 rows × 5 columns

```
In [8]: df.isnull().sum()
```

```
Out[8]: passenger_count    65441
payment_type      65441
fare_amount        0
trip_distance      0
duration          0
dtype: int64
```

```
In [9]: (65441/len(df))*100
```

```
Out[9]: 1.021716132126611
```

```
In [10]: df.dropna(inplace = True)
```

```
In [11]: df
```

```
Out[11]:   passenger_count  payment_type  fare_amount  trip_distance  duration
0             1.0           1.0         6.0          1.20       4.800000
1             1.0           1.0         7.0          1.20       7.416667
2             1.0           1.0         6.0          0.60       6.183333
3             1.0           1.0         5.5          0.80       4.850000
4             1.0           2.0         3.5          0.00       2.300000
...
6339562       1.0           1.0        11.0          2.10      14.233333
6339563       1.0           1.0        13.0          2.13      19.000000
6339564       1.0           1.0        12.5          2.55      16.283333
6339565       1.0           2.0         8.5          1.61      9.633333
6339566       1.0           1.0         0.0          0.00      1.066667
```

6339567 rows × 5 columns

```
In [12]: df['passenger_count'] = df['passenger_count'].astype('int64')
df['payment_type'] = df['payment_type'].astype('int64')
```

```
In [13]: df[df.duplicated()]
```

```
Out[13]:
```

	passenger_count	payment_type	fare_amount	trip_distance	duration
2056	1	2	7.0	0.00	0.000000
2441	1	1	52.0	0.00	0.200000
2446	2	1	9.5	1.70	13.066667
2465	1	1	4.0	0.40	3.083333
3344	1	1	6.0	1.20	5.350000
...
6339558	1	2	8.0	1.63	8.800000
6339559	1	1	8.5	1.81	8.016667
6339560	1	2	6.5	0.98	6.900000
6339562	1	1	11.0	2.10	14.233333
6339565	1	2	8.5	1.61	9.633333

3331706 rows × 5 columns

```
In [14]: df.drop_duplicates(inplace = True)
```

```
In [15]: df.shape
```

```
Out[15]: (3007861, 5)
```

```
In [16]: df['passenger_count'].value_counts(normalize = True)
```

```
Out[16]: passenger_count
1    0.581981
2    0.190350
3    0.066360
5    0.062937
6    0.039272
4    0.036046
0    0.023033
7    0.000009
9    0.000006
8    0.000006
Name: proportion, dtype: float64
```

```
In [17]: df['payment_type'].value_counts(normalize = True)
```

```
Out[17]: payment_type
1    6.782670e-01
2    3.075731e-01
3    8.721480e-03
4    5.438084e-03
5    3.324622e-07
Name: proportion, dtype: float64
```

```
In [18]: df = df[(df['passenger_count'] > 0) & (df['passenger_count'] < 6)]
df = df[df['payment_type'] < 3]
```

```
In [19]: df['payment_type'].replace([1,2],['Card','Cash'], inplace = True)
```


```
In [20]: df.describe()
```

```
Out[20]:
```

	passenger_count	fare_amount	trip_distance	duration
count	2.780283e+06	2.780283e+06	2.780283e+06	2.780283e+06
mean	1.733386e+00	1.780567e+01	4.536729e+00	2.415478e+01
std	1.176652e+00	1.506997e+01	4.895890e+00	9.260031e+01
min	1.000000e+00	-5.000000e+02	-2.218000e+01	-2.770367e+03
25%	1.000000e+00	9.000000e+00	1.500000e+00	9.883333e+00
50%	1.000000e+00	1.300000e+01	2.730000e+00	1.573333e+01
75%	2.000000e+00	2.100000e+01	5.470000e+00	2.336667e+01
max	5.000000e+00	4.265000e+03	2.628800e+02	8.525117e+03

```
In [21]: df = df[df['fare_amount'] > 0]
df = df[df['trip_distance'] > 0]
df = df[df['duration'] > 0]
```

```
In [22]: sns.boxplot(data = df, y = 'fare_amount', x = 'payment_type')
plt.show()
```



```
In [23]: for col in ['fare_amount', 'trip_distance', 'duration']:
    q1 = df[col].quantile(0.25)
    q3 = df[col].quantile(0.75)
    IQR = q3-q1

    lower_bound = q1-1.5*IQR
    upper_bound = q3+1.5*IQR

    df = df[(df[col] >= lower_bound) & (df[col] <= upper_bound)]
```

```
In [24]: df
```

Out[24]:


	passenger_count	payment_type	fare_amount	trip_distance	duration
0	1	Card	6.0	1.20	4.800000
1	1	Card	7.0	1.20	7.416667
2	1	Card	6.0	0.60	6.183333
3	1	Card	5.5	0.80	4.850000
5	1	Cash	2.5	0.03	0.883333
...
6339550	4	Card	10.5	2.40	12.383333
6339555	3	Card	10.0	2.09	14.800000
6339561	1	Card	17.5	4.11	21.500000
6339563	1	Card	13.0	2.13	19.000000
6339564	1	Card	12.5	2.55	16.283333

2297908 rows × 5 columns

In [26]:

```
plt.figure(figsize = (12,5))
plt.subplot(1,2,1)
plt.title('Distribution of fare amount')
plt.hist(df[df['payment_type'] == 'Card']['fare_amount'], histtype = 'barstacked')
plt.hist(df[df['payment_type'] == 'Cash']['fare_amount'], histtype = 'barstacked')
plt.legend()

plt.subplot(1,2,2)
plt.title('Distribution of trip distance')
plt.hist(df[df['payment_type'] == 'Card']['trip_distance'], histtype = 'barstacked')
plt.hist(df[df['payment_type'] == 'Cash']['trip_distance'], histtype = 'barstacked')
plt.legend()
plt.show()
```




```
In [27]: df.groupby('payment_type').agg({'fare_amount':['mean','std'], 'trip_distance':
```

```
Out[27]:
```

payment_type	fare_amount		trip_distance	
	mean	std	mean	std
Card	13.112493	5.849281	2.992237	1.99274
Cash	11.758005	5.613038	2.602207	1.91372

```
In [29]: plt.title('Preference of Payment Type')
plt.pie(df['payment_type'].value_counts(normalize = True), labels = df['paymen
    startangle = 90, shadow = True, autopct = '%1.1f%%', colors = ['#FA643
plt.show()
```



```
In [35]: passenger_count = df.groupby(['payment_type', 'passenger_count'])[['passenger_
passenger_count.rename(columns = {'passenger_count':'count'}, inplace = True)
passenger_count.reset_index(inplace = True)
```

```
In [39]: passenger_count['perc'] = round(passenger_count['count']/passenger_count['cour
```

```
In [40]: passenger_count
```

Out[40]:

	payment_type	passenger_count	count	perc
0	Card	1	909245	39.57
1	Card	2	327661	14.26
2	Card	3	122412	5.33
3	Card	4	63676	2.77
4	Card	5	124045	5.40
5	Cash	1	460550	20.04
6	Cash	2	155472	6.77
7	Cash	3	54506	2.37
8	Cash	4	32715	1.42
9	Cash	5	47626	2.07

In [43]:

```
df2 = pd.DataFrame(columns = ['payment_type',1,2,3,4,5])
df2['payment_type'] = ['Card','Cash']
df2.iloc[0,1:] = passenger_count.iloc[0:5,-1]
df2.iloc[1,1:] = passenger_count.iloc[5:,-1]
df2
```


Out[43]:

	payment_type	1	2	3	4	5
0	Card	39.57	14.26	5.33	2.77	5.4
1	Cash	20.04	6.77	2.37	1.42	2.07


In [48]:

```
fig, ax = plt.subplots(figsize = (20,6))
df2.plot(x = 'payment_type', kind = 'barh', stacked = True, ax = ax ,color = [ 'red','blue'])

for p in ax.patches:
    width = p.get_width()
    height = p.get_height()
    x, y = p.get_xy()
    ax.text(x + width / 2,
            y + height / 2,
            '{:.0f}%'.format(width),
            horizontalalignment = 'center',
            verticalalignment = 'center')
```



```
In [51]: fig = sm.qqplot(df['fare_amount'], line = '45')
plt.show()
```



```
In [52]: card_sample = df[df['payment_type'] == 'Card']['fare_amount']
cash_sample = df[df['payment_type'] == 'Cash']['fare_amount']
```

```
In [53]: t_stats, p_value = st.ttest_ind(a = card_sample, b = cash_sample, equal_var =
print(f'T statistic: {t_stats}, p-value: {p_value}')
```

T statistic: 169.2111527245052, p-value: 0.0

```
In [55]: if p_value < 0.05:
    print('Null hypothesis is rejected')
else:
    print('Null hypothesis is accepted')
```

Null hypothesis is rejected

```
In [72]: result = linregress(df['duration'],df['fare_amount'])
```

```
In [73]: print("Slope:", result.slope)
print("Intercept:", result.intercept)
print("R²:", result.rvalue**2)
print("P-value:", result.pvalue)
```

Slope: 0.6922301341983748
Intercept: 2.5274733774071585
R²: 0.7601694455031354
P-value: 0.0

```
In [74]: print(f"Fare Amount = {result.slope:.2f} * Duration + {result.intercept:.2f}")
```

Fare Amount = 0.69 * Duration + 2.53

```
In [75]: y_line = result.slope * df['duration'] + result.intercept

plt.scatter(df['duration'], df['fare_amount'])
plt.plot(df['duration'], y_line)
plt.xlabel("Duration")
plt.ylabel("Fare Amount")
plt.show()
```

